Step 6 — model size, search effort and memory

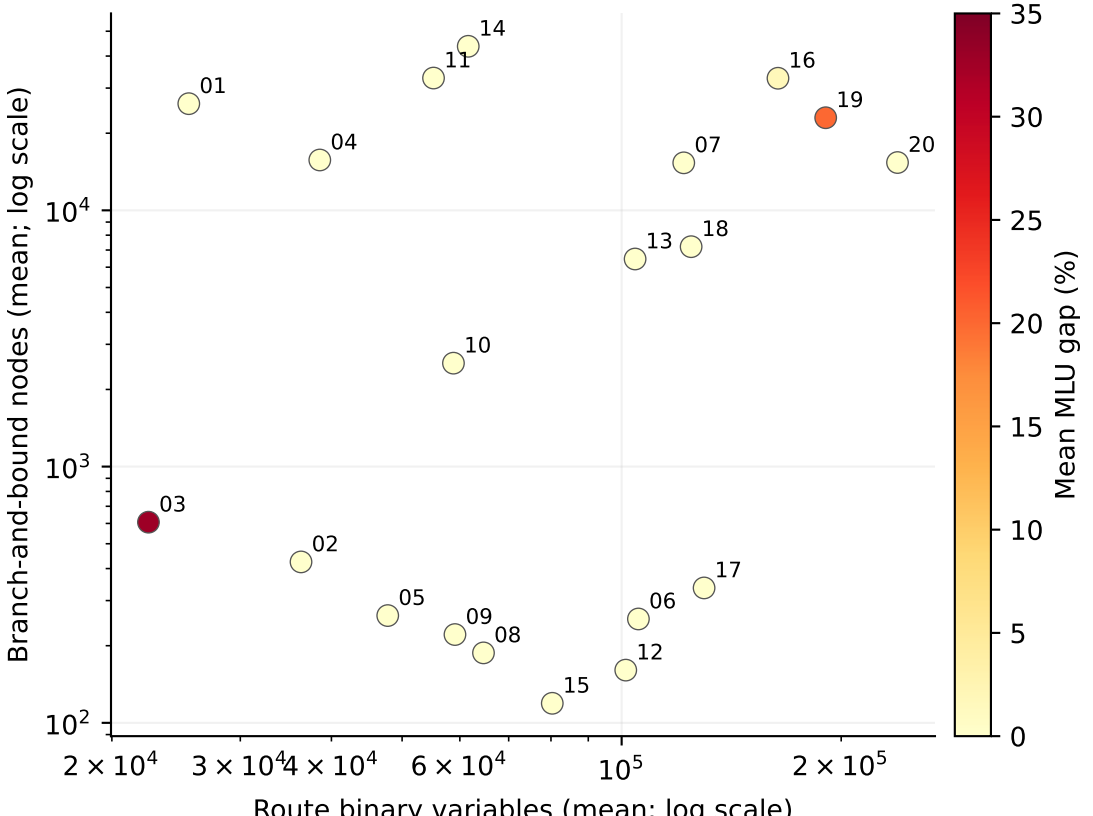

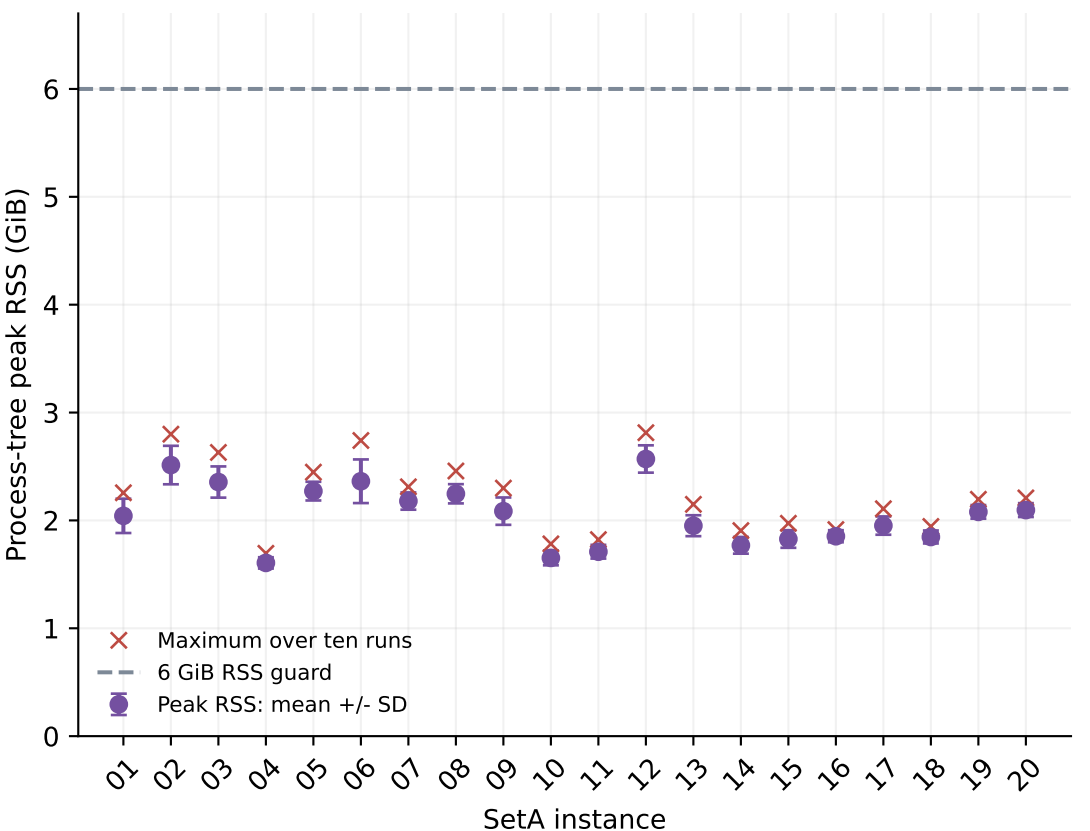

Node count and size alone do not explain solution quality: small SetA-03 still has a 32.84% MLU gap.
Maximum observed peak RSS was 2.81 GiB, below the 6 GiB per-job guard; these measurements apply to Step 6 only.